

PROJECT: Machine Learning

MNIST Classifier

Purpose of the Document

The document has to specify the requirements for the project “Build an MNIST Classifier.” Apart from specifying the functional and non-functional requirements for the project, it also serves as an input for project scoping.

Problem Statement

The purpose of the project is to train a model on the MNIST image database to detect images with 5 digits. This is done using one or more of the classification algorithms.

You are provided with the following resources that can be used as inputs for your model:

1. A collection of images of 70,000 handwritten digits is given as part of the Scikit-Learn “datasets” module. You need to import it into code.
2. Code template containing following code blocks:
 - a. Import modules (part 1)
 - b. Plot functions (part 2)
 - c. Binary classifier using SGDClassifier class (part 3)
 - d. Predict and validate using cross validation (part 4)
 - e. Print Confusion Matrix, precision, and recall (part 5)

You are expected to write the code for a binary classification model using Python Scikit-Learn that trains on data and calculates the accuracy score on the test data.

Project Guidelines

#	Exercises	Process
1	Initiation	Begin by extracting the ipynb file.
2	Exercise 1	<p>Build a 5 vs not-5 binary classifier using KNN with number of neighbors as 4.</p> <p>(Hint: Use Scikit-Learn library KNeighboursClassifier)</p> <p>(Hint: Refer to the KNN tutorial taught earlier in the course)</p> <p>Print accuracy score of the model.</p> <p>Note that it might take 10 minutes to print the accuracy score.</p>
3	Exercise 2	<p>Build a 5 vs not-5 binary classifier using Logistic Regression.</p> <p>(Hint: Use Scikit-Learn library LogisticRegression)</p> <p>(Hint: Refer to the logistic regression tutorial taught earlier in the course)</p> <p>Print confusion matrix, accuracy score, and cross validation score.</p> <p>Note that it might take 10 minutes to print the accuracy score.</p>
4	Exercise 3	<p>Build a 5 vs not-5 binary classifier using SVMs</p> <p>(Hint: Refer to the SVM tutorial taught earlier in the course)</p> <p>Print accuracy score.</p> <p>Note that it might take 10 minutes to print the accuracy score.</p> <p>To make the processing faster, you could work with only 10000 of the samples, not the entire 70000 images.</p> <p>(Hint: <code>X_train, X_test, y_train, y_test = X[30000:37000], X[37000:40000], y[30000:37000], y[37000:40000]</code>)</p> <p>(Hint: <code>shuffle_index = np.random.permutation(7000)</code>)</p>

